


# Recitation 8-11/04/2022

## Overfitting & Regularization


TA: Akshay Antony

# Underfitting

- Underfitting is when the model is **unable to fit well** to the training data.
  - This happens when the model is very simple
  - We lose accuracy.
  - Also if training data is very low
  - Known as Bias




Regression




Classification

# Overfitting

- Overfitting is when the model **over-fits** to the training data.
  - This happens when we choose a complex model even when a simpler model can infer the ground truth well
  - We lose generalizability.
  - Also if training data is very specific.
  - Known as Variance




Regression



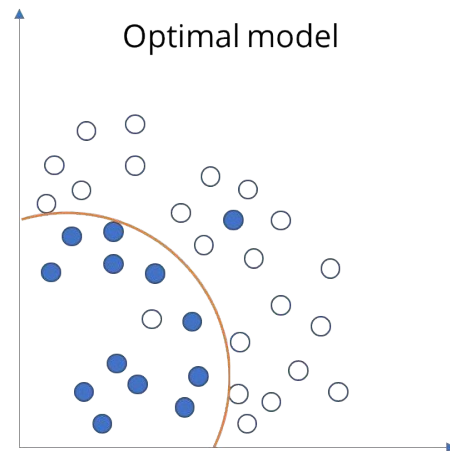
Classification

# Just Right

- Just right/optimal is when the model neither **over** nor **under** fits.

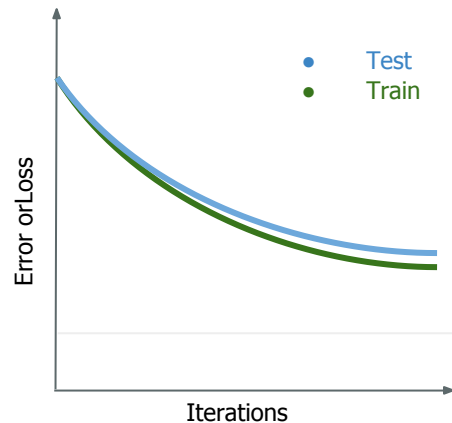


Regression

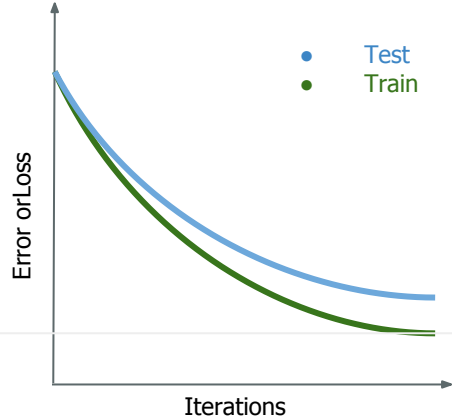


Classification


# How to visualize during training/testing



Underfitting




Just  
Right



Overfitting

# Complexity of model



# Regularization (to prevent overfitting)

Let's assume:

$$H = \theta_1 * X_1 + \theta_2 * X_2$$

If we are overfitting, it is likely because the model is giving really high importance to features, some of which might not be even useful.

So if we can reduce  $\theta$  (coefficients of the features) we can reduce the effect each feature has on the output.

Thus regularization is a way of capping the weights so they don't grow too much.

# L1 Regularization

Lasso Regression (Least Absolute Shrinkage and Selection Operator) adds “absolute value of magnitude” of coefficient as penalty term to the loss function.

- Expression:  $\sum_{i=1}^n (Y_i - \sum_{j=1}^p X_{ij}\beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j|$
- if  $\lambda=0$  reduces to unregularized case
- Can make weights go to 0 (derivative does not contain the weight term)
- More expensive
- Not closed form solution




# L2 Regularization

Ridge regression adds “squared magnitude” of coefficient as penalty term to the

$$\sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2.$$


- Expression:
- Cannot make weights to zero (derivative depends on the weight)
- Less expensive
- Closed form (differentiable at every point)




**FIGURE 3.11.** Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions  $|\beta_1| + |\beta_2| \leq t$  and  $\beta_1^2 + \beta_2^2 \leq t^2$ , respectively, while the red ellipses are the contours of the least squares error function.

# Dropout

1. Dropout is a regularization method that approximates training a large number of neural networks with different architectures in parallel.
2. Dropout has the effect of making the training process noisy, forcing nodes within a layer to probabilistically take on more or less responsibility for the inputs.
3. Dropout simulates a sparse activation from a given layer, which interestingly, in turn, encourages the network to actually learn a sparse representation as a side-effect



(a) Standard Neural Net



(b) After applying dropout.

# Pytorch nn.Dropout


- `torch.nn.Dropout(p=0.5, inplace=False)`
- `p` (float) – probability of an element to be zeroed. Default: 0.5
- Scaled by during training factor of  $1/(1-p)$
- If `model.eval()` dropout is deactivated

Where Not to use:


- Just before final linear layer
- For small networks
- Not training for large number of iterations

# Batchnorm as regularizer

- Normalizing the inputs to the layer has an effect on the training of the model, dramatically reducing the number of epochs required.
- In pytorch BatchNorm learns a mean and std for each BatchNorm layer
- During training this mean and std acts as noise and hence a regularizer
- In eval mode the learned mean and std is used



Before standardization




After standardization

# Data Augmentation

Images (torchvision.transforms)

## 1. Simple transformations

- Resize: `Resize(size)`
- Gray Scale: `Grayscale()`, `to_grayscale()`
- Normalize: `Normalize(mean, std)`
- Random Rotation: `RandomRotation(degree)`
- Center Crop: `CenterCrop(size)`
- Random Crop: `RandomCrop(size)`



# Random Rotation

- `torchvision.transforms.RandomRotation(degrees)`
- degrees (sequence or number) – Range of degrees to select from



# GrayScale

- `torchvision.transforms.Grayscale`

Original image




# Normalize

- `torchvision.transforms.Normalize(mean, std)`
- mean: mean of size (C, 1) C number of channels in input image
- std: std of size (C, 1) C number of channels in input image
- mean = [0.485, 0.456, 0.406] and std = [0.229, 0.224, 0.225] for imagenet, use this while doing transfer learning

# RandomResize

- `torchvision.transforms.RandomResizedCrop`
- Crop a random portion of image and resize it to a given size.

Original image



# Flips

- RandomHorizontalFlip: argument probability to flip
- RandomVerticalFlip: argument probability to flip

Original image



Original image



## Custom Augmentation

- Write a class capable of using transforms.Compose
- Define a function named `__call__` (object() is shorthand for object.\_\_call\_\_())


```
class RandomHorizontalFlip(object):  
    def __init__(self, p=0.5):  
        self.p = p  
  
    def __call__(self, sample):  
        image, target = sample['image'],  
            sample['target']  
  
        if random.random() > self.p:  
            trans = tt.RandomHorizontalFlip(1)  
            image = trans(image)  
  
        if target != 0:  
            target = -1 * target  
  
        return {'image': image, 'target': target}
```



# Data Augmentation

## Point Clouds:

- Write custom transforms
- Random Rotation
- Random flipping
- Adding Noise
- Random scaling
- Drop Points




# Early Stopping

- When training a large network, there will be a point during training when the model will stop generalizing and start learning the statistical noise in the training dataset.
- If the performance of the model on the validation dataset starts to degrade (e.g. loss begins to increase or accuracy begins to decrease), then the training process is stopped.
- Early stopping may be thought of as a type of “implicit” regularization, much like using a smaller network that has less capacity.

## Criteria:

- $\text{val\_loss} - \text{train\_loss} > \text{threshold}$
- $\text{val\_loss}$  does not improve



Available in pytorch-lightning

```
EarlyStopping(monitor="val_accuracy",  
min_delta=0.00, patience=3,  
verbose=False, mode="max")
```

1. monitor: loss needed to focus on
2. min\_delta: amount of change to consider as an improvement
3. patience: number iteration to wait until stopping